

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import RobustScaler, OneHotEncoder
from sklearn.linear_model import LogisticRegression
```

```
In [4]: df = pd.read_csv('churn modelling.csv', index_col=0)
df.head()
```

Out[4]:

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
RowNumber													
1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [5]: df.drop(['CustomerId', 'Surname'], axis=1, inplace=True)
```

```
In [6]: df.shape
```

Out[6]: (10000, 11)

```
In [7]: df.isna().sum()
```

Out[7]: CreditScore 0  
Geography 0  
Gender 0  
Age 0  
Tenure 0  
Balance 0  
NumOfProducts 0  
HasCrCard 0  
IsActiveMember 0  
EstimatedSalary 0  
Exited 0  
dtype: int64

```
In [8]: X = df.drop('Exited', 1)
y = df.Exited
y.value_counts()
```

Out[8]: 0 7963  
1 2037  
Name: Exited, dtype: int64

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0, stratify=y)
X.columns
```

Out[9]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',  
 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary'],  
 dtype='object')

```
In [10]: num_cols = ['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
cat_cols = ['HasCrCard', 'IsActiveMember', 'Geography', 'Gender']
```

```
In [11]: ct = ColumnTransformer([
            ('s1', RobustScaler(), num_cols),
            ('s2', OneHotEncoder(sparse=False, handle_unknown='ignore'), cat_cols)
        ])
```

```
In [12]: p = Pipeline([
            ('ct', ct),
            ('mod', LogisticRegression(random_state=0))
        ])
```

```
In [13]: p.fit(X_train, y_train)
```

```
Out[13]: Pipeline(steps=[('ct',
                           ColumnTransformer(transformers=[('s1', RobustScaler(),
                                                             ['CreditScore', 'Age',
                                                             'Tenure', 'Balance',
                                                             'NumOfProducts',
                                                             'EstimatedSalary']),
                                                           ('s2',
                                                            OneHotEncoder(handle_unknown='ignore',
                                                                           sparse=False),
                                                             ['HasCrCard',
                                                             'IsActiveMember',
                                                             'Geography', 'Gender'])])),
                          ('mod', LogisticRegression(random_state=0))])
```

```
In [14]: preds = p.predict(X_test)
preds[:15]
```

Out[14]: array([1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=int64)

```
In [15]: np.array(y_test)[:15]
```

Out[15]: array([1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0], dtype=int64)

```
In [16]: from sklearn.metrics import confusion_matrix, plot_confusion_matrix
confusion_matrix(y_true=y_test, y_pred=preds)
```

Out[16]: array([[1530, 63],
 [ 319, 88]], dtype=int64)

```
In [17]: p.classes_
```

Out[17]: array([0, 1], dtype=int64)

```
In [18]: confusion_matrix(y_test, preds, labels=(1,0))
```

Out[18]: array([[ 88, 319],
 [ 63, 1530]], dtype=int64)

```
In [19]: confusion_matrix(y_test, preds, labels=(1,0)).ravel()
```

Out[19]: array([ 88, 319, 63, 1530], dtype=int64)

```
In [38]: accuracy_score(y_test, preds)
```

Out[38]: 0.809

```
In [39]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,\
fbeta_score, matthews_corrcoef
precision_score(y_test, preds)
```

Out[39]: 0.5827814569536424

```
In [40]: tp, fn, fp, tn = confusion_matrix(y_test, preds, labels=(1,0)).ravel()
precision = tp/(tp+fp)
precision
```

Out[40]: 0.5827814569536424

```
In [41]: recall_score(y_test, preds)
```

Out[41]: 0.21621621621621623

```
In [42]: # harmonic mean of precision and recall
f1_score(y_test, preds)
```

Out[42]: 0.31541218637992835

```
In [ ]:
```

```
In [ ]:
```